

Database Models: Beyond RDBMS

UA.DETI.CBD
José Luis Oliveira / Carlos Costa

The Battle of the Data Models

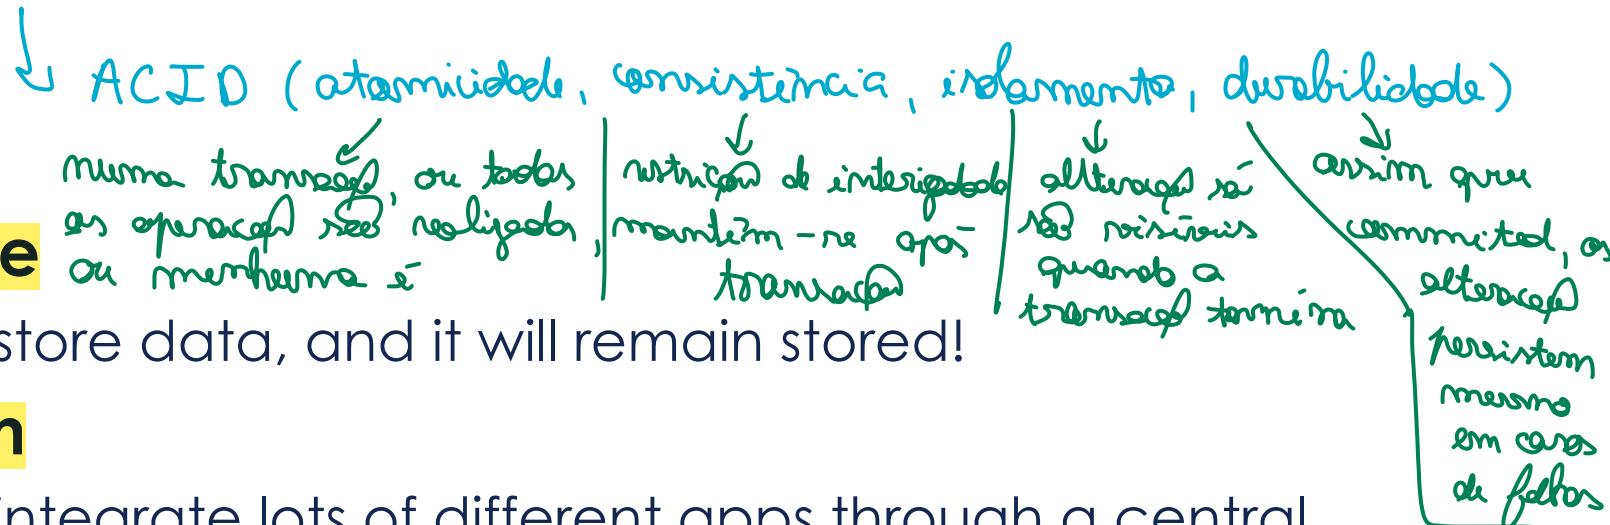
- ❖ Data models are perhaps the most important part of developing software

(os modelos de dados (como quais os programas no database) têm um papel fundamental na sua programação)

When we have some data...

❖ Relational Databases solve most data problems

Why?



❖ Integration

- We can integrate lots of different apps through a central DB

❖ SQL

- Standard, well understood, very expressive

❖ Transactions

- ACID transactions, strong consistency

Transactions – ACID Properties

❖ Atomic

- All of the work in a transaction completes (commit) or none of it completes

❖ Consistent

- A transaction transforms the database from one consistent state to another consistent state. Consistency is defined in terms of constraints.

❖ Isolated

- The results of any changes made during a transaction are not visible until the transaction has committed. Concurrent interactions behave as though they occurred serially

❖ Durable

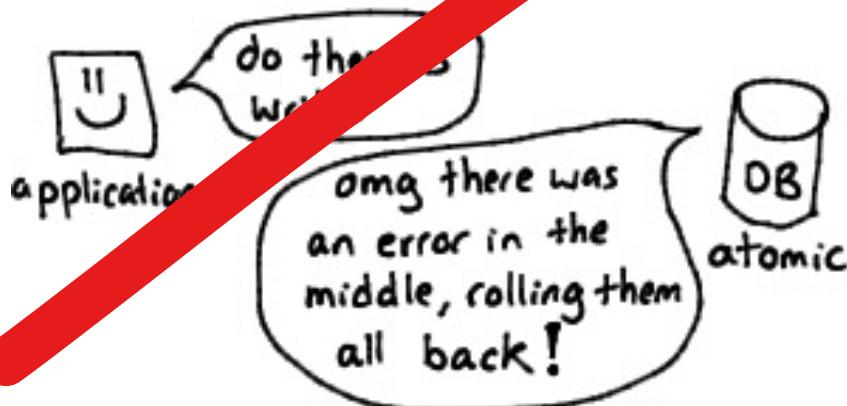
- The results of a committed transaction survive failures

ACID

ACID is about safety guarantees for database transactions.



not about concurrent writes
(that's "isolation")

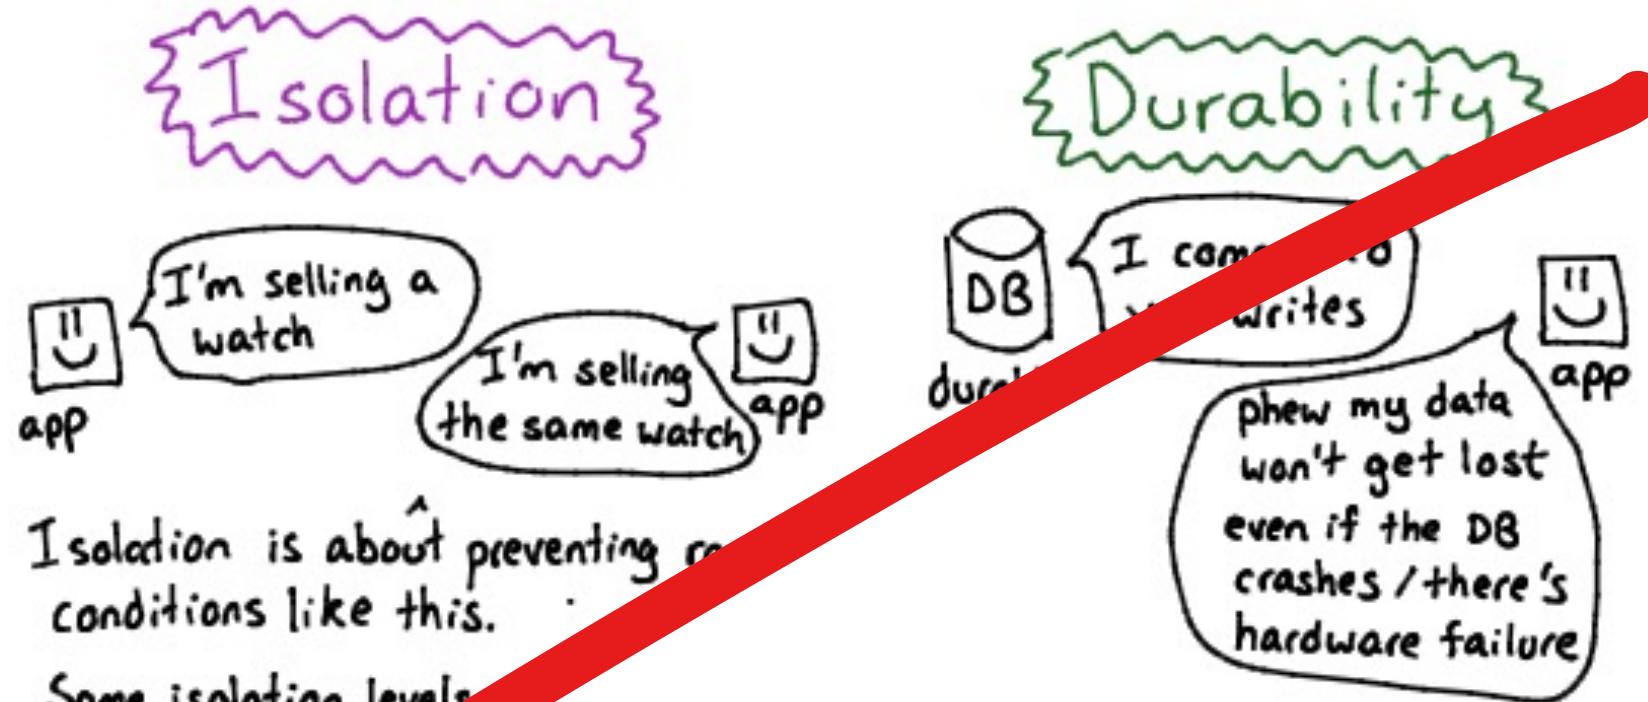


super overloaded term.
This sense of "consistency" is actually an application property not a DB property.

not serializability
not as in eventual consistency

About preserving application invariants like "every user gets an invoice".

ACID



Isolation is about preventing conditions like this.

Some isolation levels

- serializability
- snapshot isolation
- read committed

Perfect durability doesn't exist.

Can involve:

- write-ahead log (usually)
- replication

<https://jvns.ca/blog/2017/06/11/log-structured-storage/>

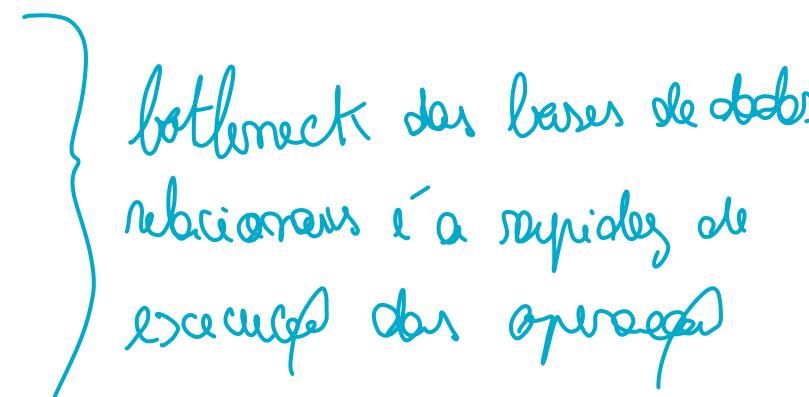
The Relational Model

- ❖ The relational model, proposed by Edgar Codd in 1970, is still the best-known data model today.
 - data is organized into relations (in SQL: tables), where each relation is an unordered collection of tuples (rows).
 - The goal of the relational model was to hide that implementation detail behind a cleaner interface.

Rivals of the Relational Model

- ❖ Over the years, there have been many competing approaches to data storage and querying.
 - Object databases came and went again in the late 1980s and early 1990s.
 - XML databases appeared in the early 2000s, but have only seen niche adoption.
- ❖ Much of what you see on the web today is still powered by relational databases
 - Online publishing, discussion, social networking, e-commerce, games, software-as-a-service productivity applications, or much more.
- ❖ Now NoSQL is the most recent attempt to overthrow the relational model's dominance

Current trends and Issues

- ❖ A few key trends and issues have motivated change in relational data storage technologies
 - ...In use cases
 - ...In technology
 - ❖ Key trends include:
 - Increasing **volume of data and traffic**
 - More complex data connectedness
 - ❖ Key Issues include:
 - The **impedance mismatch** problem
- 
- bottleneck das bases de dados
reduzir o tempo é a rapidez de
execução das operações

Impedance Mismatch

} paradigma entre principíos de engenharia de software & princípios abstracionais

❖ Object Orientation

- based on software engineering principles

conflictos

❖ Relational Paradigms

- based on mathematics and set theory

Ex.: vários objetos que representam funcionários numa empresa

• cada funcionário tem o seu departamento } mas vários funcionários podem trabalhar no mesmo

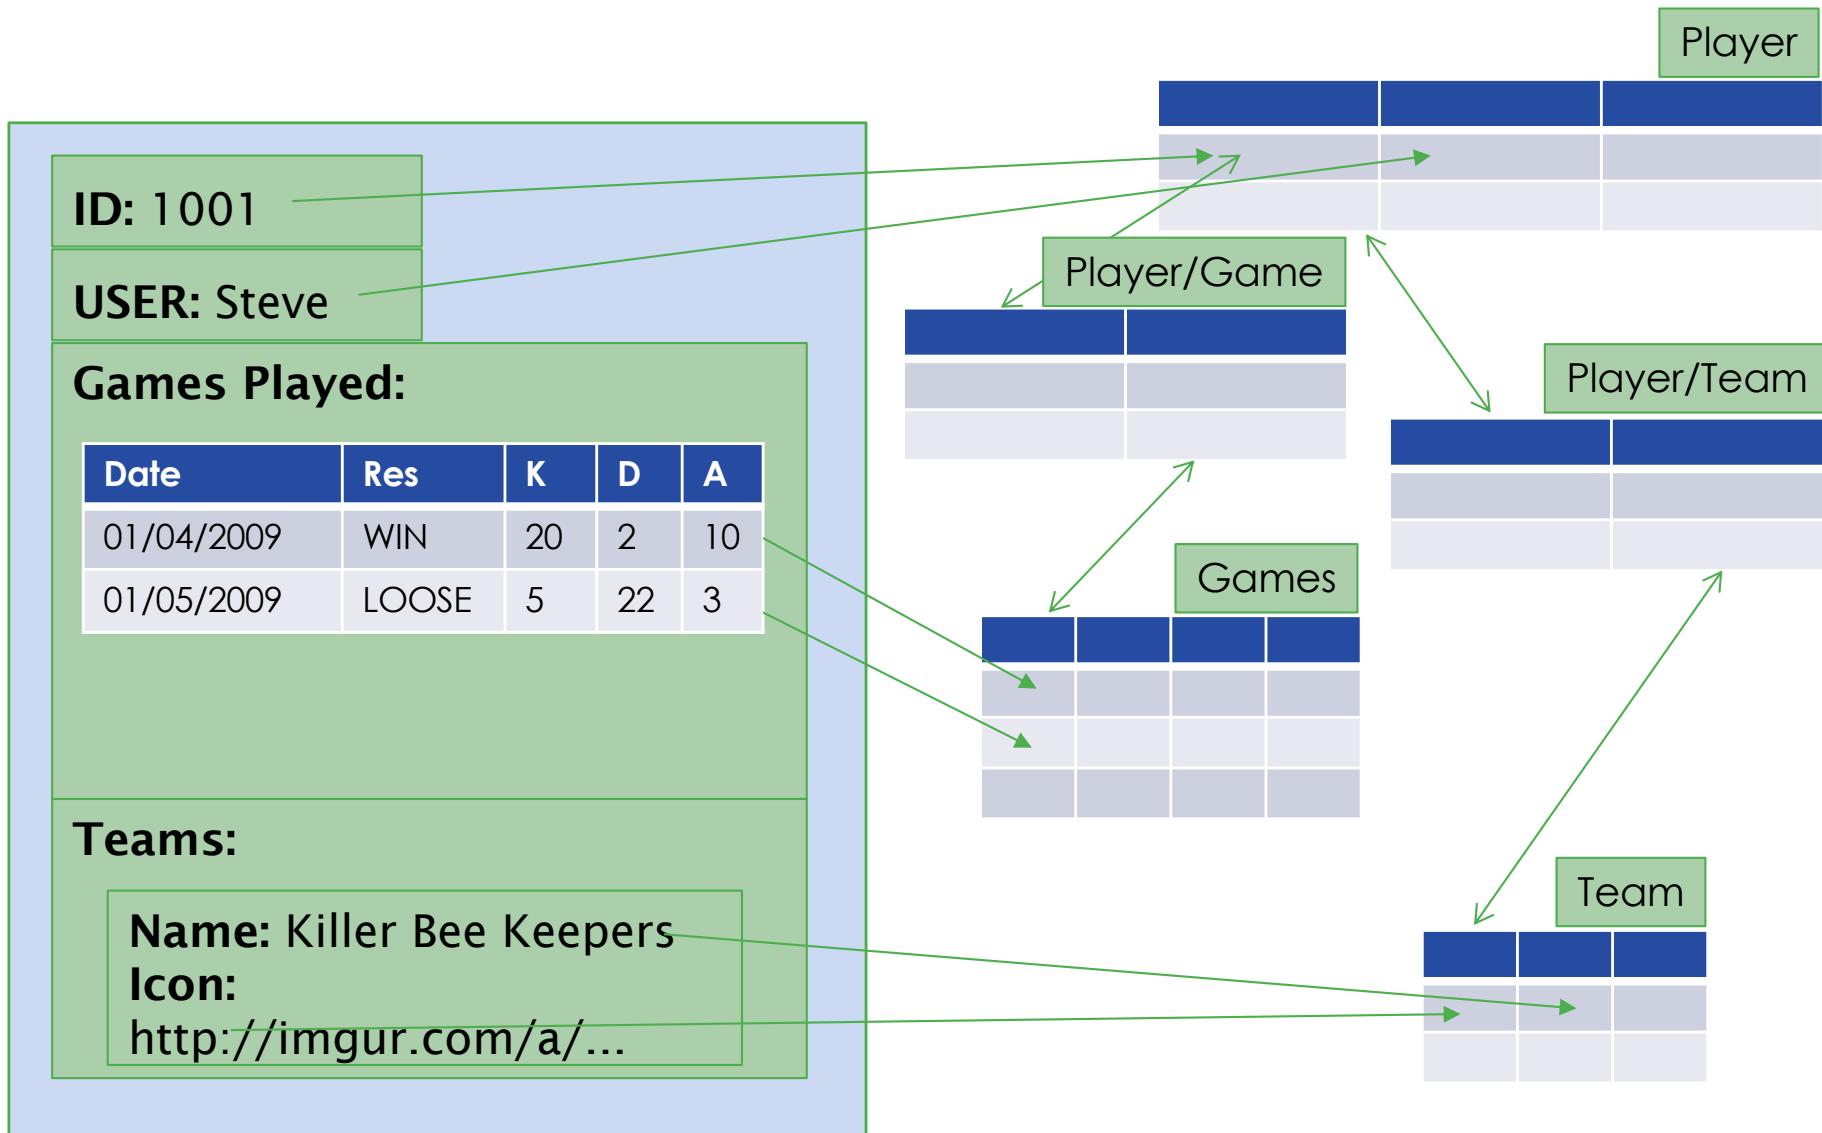
❖ To store data persistently in modern programs a single logical structure must be split up

- The nice word is **normalised**

!!

isso quer dizer
uma repetição de departamentos nos funcionários e a base de dados não estaria normalizada
(redundância de dados)

Impedance Mismatch – example



Impedance Mismatch – example

<http://www.linkedin.com/in/williamhgates>



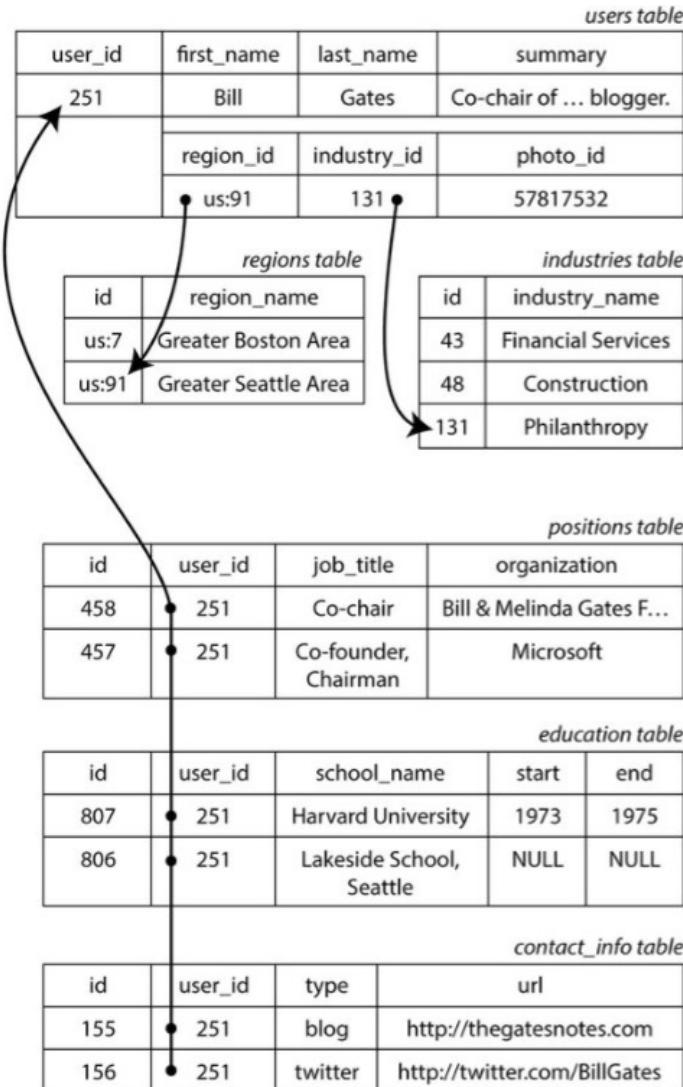
Bill Gates
Greater Seattle Area | Philanthropy

Summary
Co-chair of the Bill & Melinda Gates Foundation. Chairman, Microsoft Corporation. Voracious reader. Avid traveler. Active blogger.

Experience
Co-chair • Bill & Melinda Gates Foundation
2000 – Present
Co-founder, Chairman • Microsoft
1975 – Present

Education
Harvard University
1973 – 1975
Lakeside School, Seattle

Contact Info
Blog: thegatesnotes.com
Twitter: @BillGates



Normalization

→ visa reduzir a redundância de dados

UTILIZA IDs
(para identificar entidades)

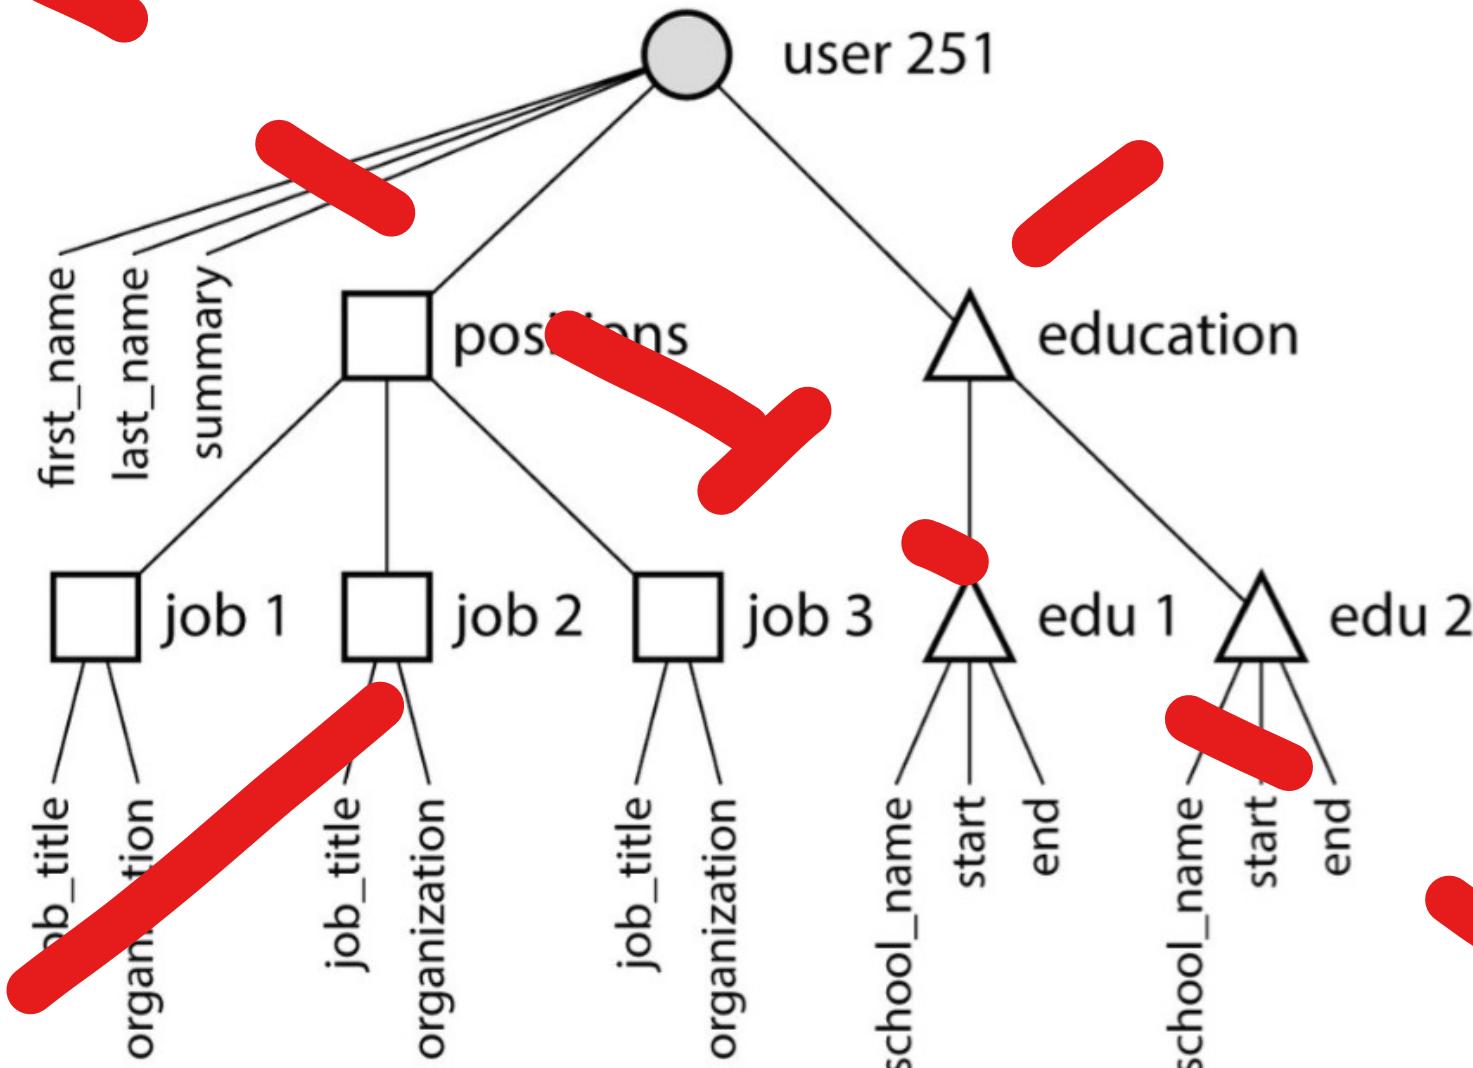
- ❖ Why IDs (region_id, industry_id, ..) and not plain-text?

- Consistent style and spelling across profiles,
- Avoiding ambiguity, e.g. if there are several cities with the same name, (*reduzem entidades similares*)
- The name is stored only in one place, so it is easy to update, (*facilita alterar - informaçõe organizada em só 1 tabela*)
- Simplify translation into other languages,

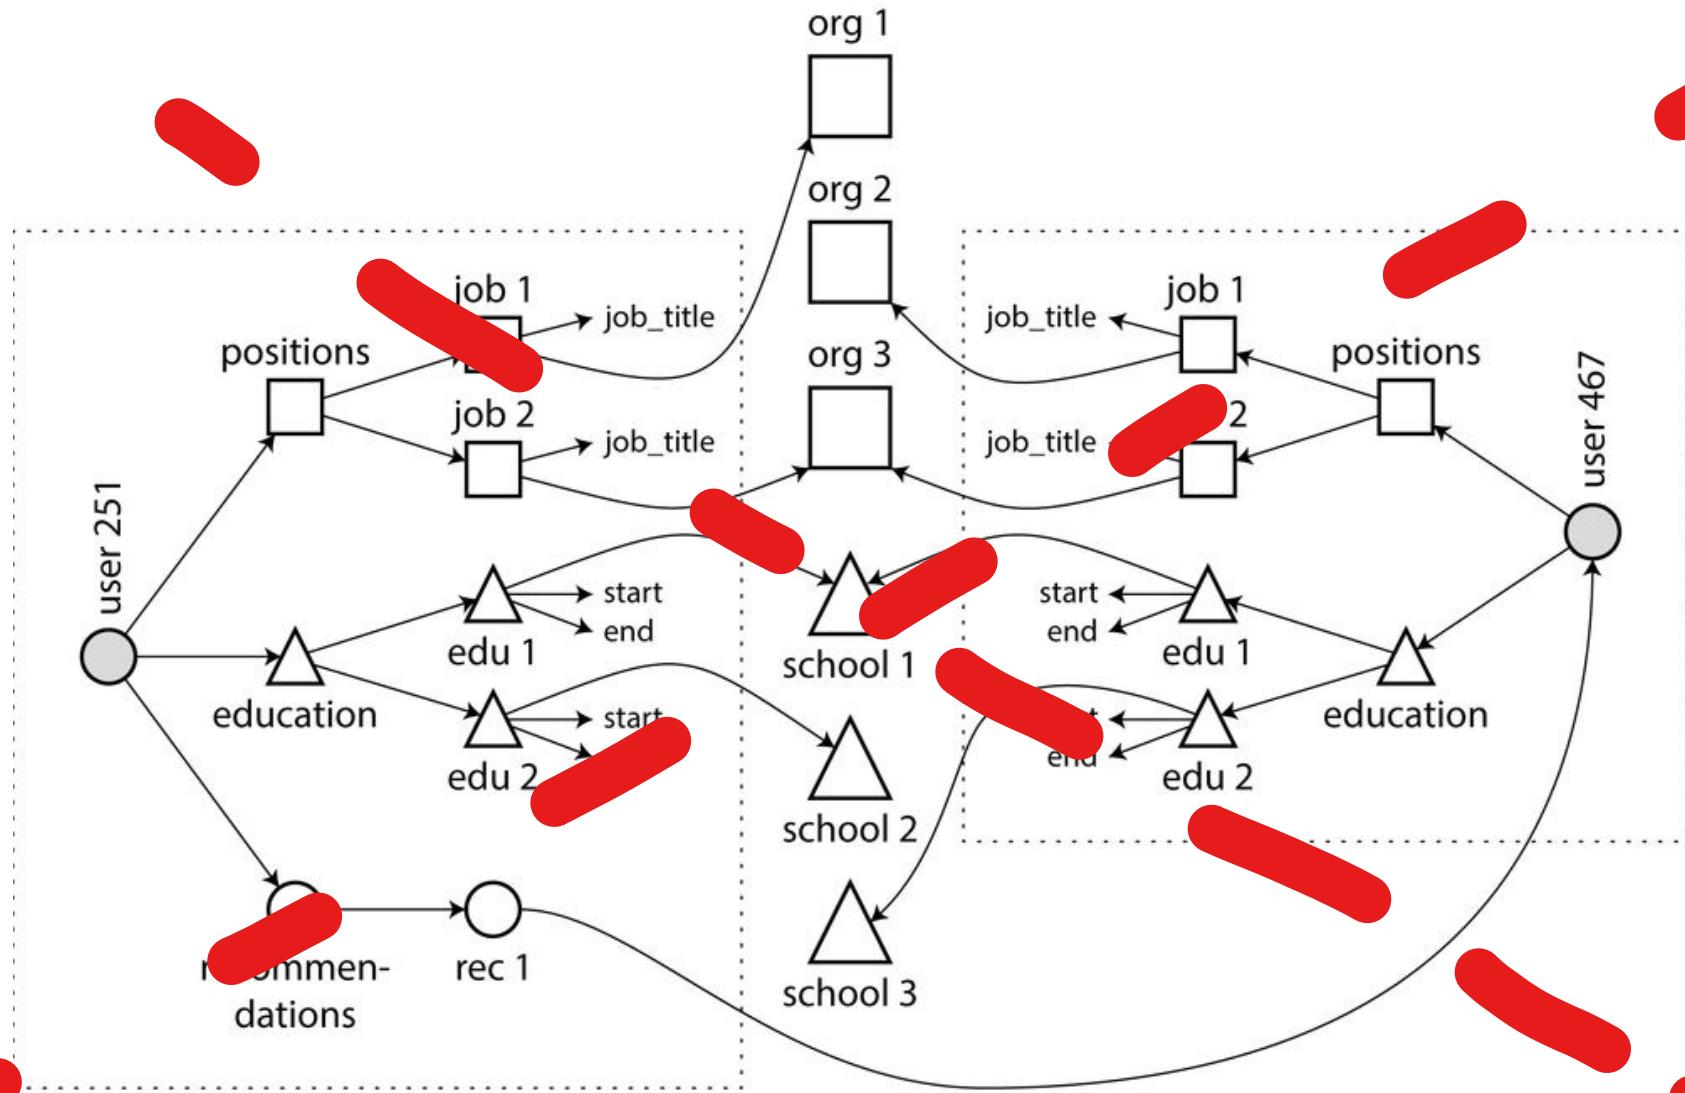
- ❖ A database in which entities like region and industry are referred to by ID is called **normalized**.

- ❖ A database that duplicates the names and properties of entities on each document is **denormalized**.

One-to-Many relations



Many-to-Many relationships



Increased Data Volume

- ❖ We are creating, storing, processing more data than ever before!
 - "From 2005 to 2020, the digital universe will grow by a factor of 300, from 130 exabytes to 40,000 exabytes, or 40 trillion gigabytes (more than 3,200 gigabytes for every man, woman, and child in 2020)."
 - THE DIGITAL UNIVERSE BY 2020: Big Data, Bigger Digital Shadows, and Biggest Growth in the Far East, Dec 2012, John Gantz and David Reinsel
 - "IDC predicts that the collective sum of the world's data will grow from 33 zettabytes this year to a 175ZB by 2025."
 - The Digitization of the World, From Edge to Core, Nov 2018
 - <https://www.seagate.com/files/www-content/our-story/trends/files/idc-seagate-dataage-whitepaper.pdf>

$$\text{EB} = 10^{18} \quad \text{ZB} = 10^{21}$$

~~Increased Data Connectivity~~

-
- ❖ The data we're producing has fundamentally changed
 - from isolated Text Documents (early 1990s)
 - ... to html pages with links (early web)
 - ... to blogs with pingback, RSS feeds (web 2.0)
 - ... to social networks (... add links between people)
 - ... to massive linked open data sets (web 3.0... one of them anyway)

Dealing with data size Trends

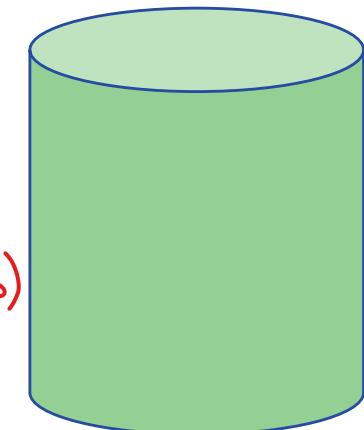
- ❖ Two options when dealing with these trends:

(construir bases de dados maiores)

- ❖ Build Bigger Database machines

- This can be expensive

- Fundamental limits to machine size (*limitações físicas*)



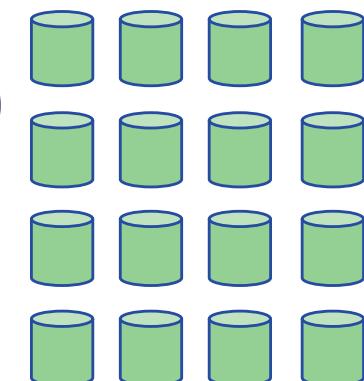
(criar um sistema distribuído de BD)

- ❖ Build Clusters of smaller machines

- Lots of small machines (commodity machines)

- Each machine is cheap, potentially unreliable

- Needs a DBMS which understands clusters



mais barata, mas menos fiável

RDBMS have fundamental issues

(Relational database management system)

❖ In dealing with (horizontal) scale

- Designed to work on single, large machines
- Difficult to distribute effectively

(bom escalamento vertical,
mais escalamento horizontal)

More subtle: An Impedance Mismatch

⚠️ Capacidade ≠ Desempenho

Escalar verticalmente (scale up) significa adicionar recursos em um único nó do sistema (mais memória ou um disco rígido mais rápido). Escalar horizontalmente (scale out) significa adicionar mais nós ao sistema, tais como um novo computador



The NoSQL Movement

(Not only SQL)

NoSQL

- ❖ The term NoSQL is unfortunate, since it doesn't refer to any technology
 - “Not only SQL”
- ❖ Nevertheless, the term struck a nerve, and quickly spread through the Web startup community and beyond.
- ❖ Several interesting database systems are now associated with the #NoSQL hashtag.

The NoSQL movement (base de dados não relacionais)

- ❖ Key attributes include:

- **Non-Relational**
 - They can be, but aren't good at it
- **Simple API**
 - **No Join**
- **BASE & CAP Theorem**
 - **No ACID requirements**
- **Schema-free**
 - Implicit schema, application side (*numa superfície de BD*)
- **Inherently Distributed**
- **Open Source**
 - mostly

BASE Transactions

(opposite to principle ACID)

em resposta às
limitações de consistência
que um banco de
dados pode ter

❖ Acronym contrived to be the opposite of ACID

- Basic Availability

- The database appears to work most of the time.

- Soft-state

- Stores don't have to be write-consistent, nor do different replicas have to be mutually consistent all the time. (escritas num só da base de dados não têm de ser escritas garantidamente em simultâneo nos outros nós).

- Eventual consistency

- Stores exhibit consistency at some later point (e.g., lazily at read time).

❖ Characteristics

- Optimistic

- Simpler and faster

- Availability first

- Best effort

- Approximate answers OK

(consistência atingida em algum momento posterior)
(quando nós são deserializados em relação a outros a determinado momento)

torna a BD + rápida

Brewer's CAP Theorem

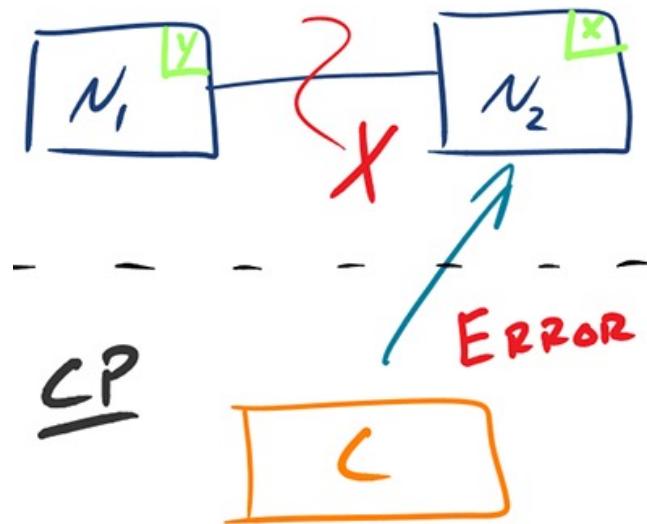
- ❖ A distributed system can support only two of the following characteristics:
 - **Consistent**
 - writes are atomic, all subsequent requests retrieve the new value
 - **Available**
 - The database will always return a value so long as the server is running
 - **Partition Tolerant** (*também é falso - um sistema consegue funcionar mesmo que 1 má deixa de responder*)
 - The system will still function even if the cluster network is partitioned (i.e. the cluster loses contact with parts of itself)
- ❖ The overly stated well cited issue is:
 - We can only ever build an algorithm which satisfies 2 of 3.
 - But .. horizontal scaling strategy is based on data partitioning;
 - Therefore, designers are forced to decide between consistency and availability.

↳ bei sempre data partitioning (P)

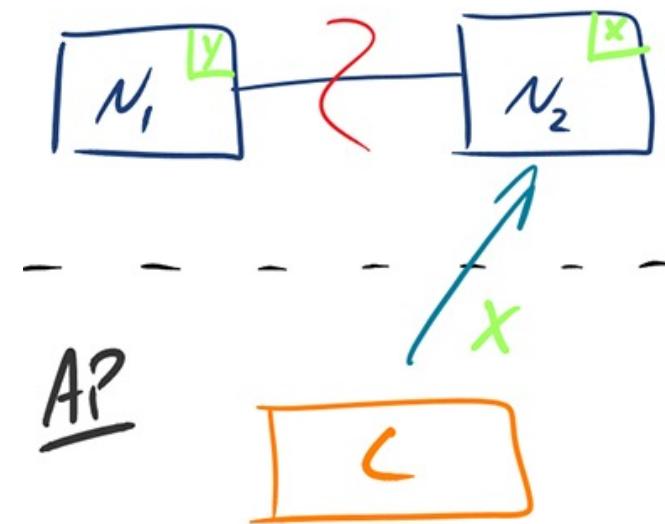
$\xrightarrow{\quad}$ CP : all clients will have the same view of the data
 $\xleftarrow{\quad}$ AP : each client can always read and write

Brewer's CAP Theorem

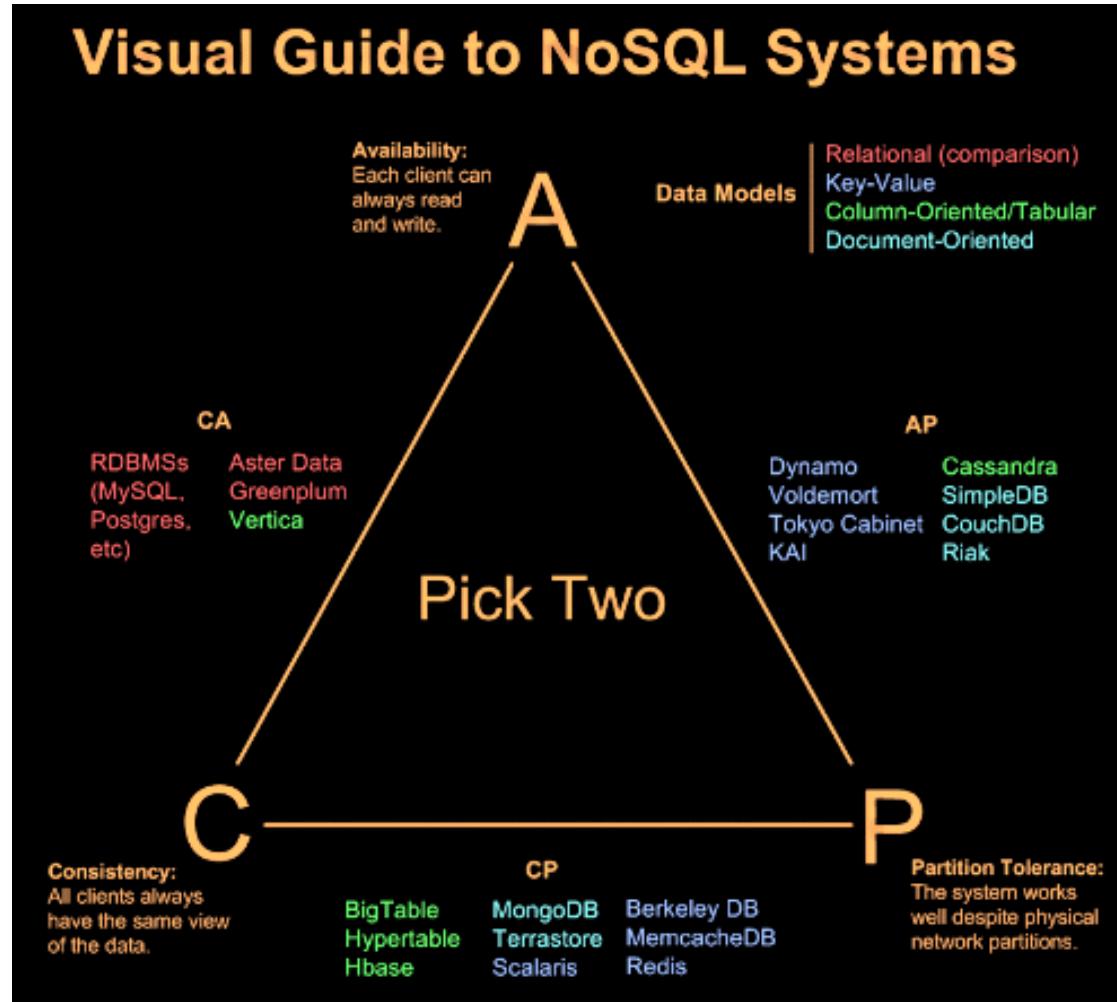
- ❖ **CP** - Consistency/Partition Tolerance



- ❖ **AP** - Availability/Partition Tolerance



CAP Theorem



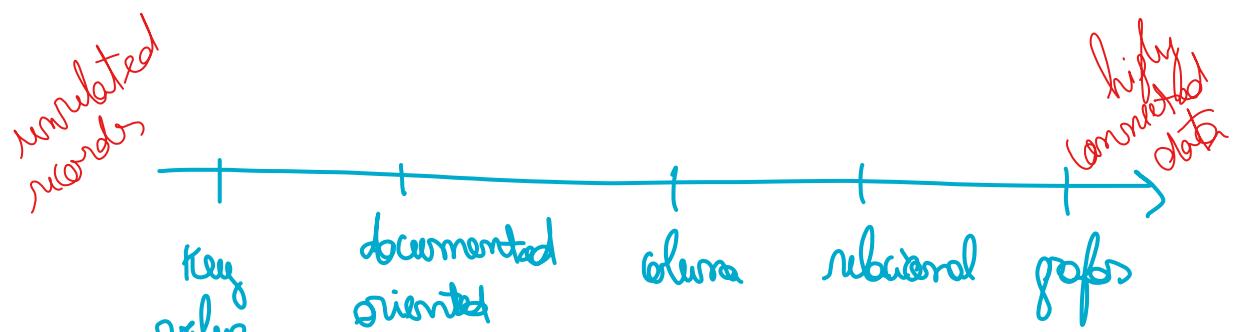
Types of NoSQL Databases

❖ Core types

- **Key-value** stores
- **Document** stores
- **Column** stores
- **Graph** databases

❖ Non-core types

- **Object** databases
- Native **XML** databases
- **RDF** stores
- ...



Key-Value Databases – Basics

❖ Data model

- The most simple NoSQL database type
- Works as a simple hash table (mapping)

❖ Key-value pairs

- **Key** (id, identifier, primary key) – usually a string.
- **Value:** can be anything (text, structure, image, etc.) – a black box for the database system.

❖ Query patterns

- Create, update or remove value for a given key
- Get value for a given key

❖ Characteristics

- great performance, easily scaled, ...
- not for complex queries nor complex data

Key-Value Databases – Basics

❖ Suitable use cases

- Session data, user profiles, user preferences, shopping carts, ...
 - I.e. when values are only accessed via keys

❖ When not to use

- Relationships among entities
- Queries requiring access to the content of the value part

Armazenam dados mais persistentes e não permitem relação entre os dados



Document Databases – Basics

estruturas compactas

é bom quando usado em

relacionamentos 1:N e raro

em N:M

o nome dos
atributos des-
creve-se a
mí própria

❖ Data model - Documents

- Self-describing complex data structure
- Hierarchical tree structures (JSON, XML, ...)
 - Scalar values, maps, lists, sets, nested documents, ...
- Identified by a unique identifier (key, ...)

❖ Document data stores understand their documents

- Queries can run against values of document fields (?)
- Indexes can be constructed for document fields

❖ Query patterns

- Create, update or remove a document
- Retrieve documents according to complex queries

❖ Difference from Key-Value stores

- Extended key-value stores. The value part is examinable!

Document Databases – Basics

```
{  
  "_id": "1",  
  "name": "steve",  
  "games_owned": [  
    {"name": "Super Meat Boy"},  
    {"name": "FTL"},  
  ],  
}
```

```
{  
  "_id": "2",  
  "name": "darren",  
  "handle": "zerocool",  
  "games_owned": [  
    {"name": "FTL"},  
    {"name": "Assassin's Creed 3", "dev": "ubisoft"},  
  ],  
}
```

JSON é um formato de troca de informações feito através de lista de dados de tipos diferentes orientado, com **dados quantitáveis** de **dados estruturados**



apenas se cobraam
as dados que
intervinem

Document Databases – Basics

❖ Suitable use cases

- Event logging, content management systems, blogs, web analytics, e-commerce applications, ...
- I.e. for **structured documents with similar schema**

❖ When not to use

- Set operations involving multiple documents
- Design of document structure is constantly changing
 - I.e. when the required level of granularity would outbalance the advantages of aggregates

⊕ simples código fácil, flexibilidade de armazenamento e melhor performance para ~~as operações~~
SIMPLES, sem JOINS (não existem)

⊖ não geram de alterar valores frequentemente

Column Databases – Basics

❖ Data model

- Column family (table)
 - Table is a collection of similar rows (not necessarily identical)
- Row
 - Row is a collection of columns - should encompass a group of data that is accessed together
 - Associated with a unique row key
- Column *→ um termo de pesquisa, não passamos de 1 ou 2 tipos todos, só a coluna*
 - Column consists of a column name and column value (and possibly other metadata records)
 - Scalar values, but also flat sets, lists or maps may be allowed

os dados são armazenados em colunas
em vez de linhas



❖ Query patterns

- Create, update or remove a row within a given column family
- Select rows according to a row key or simple conditions

Column Databases – Basics

❖ Suitable use cases

- Event logging, content management systems, blogs, ...
 - i.e. for structured flat data with similar schema
- Batch processing via mapreduce

❖ When not to use

- ACID transactions are required
- Complex queries: aggregation (SUM, AVG, ...), joining, ...
- Early prototypes: i.e. when database design may change

❖ Examples

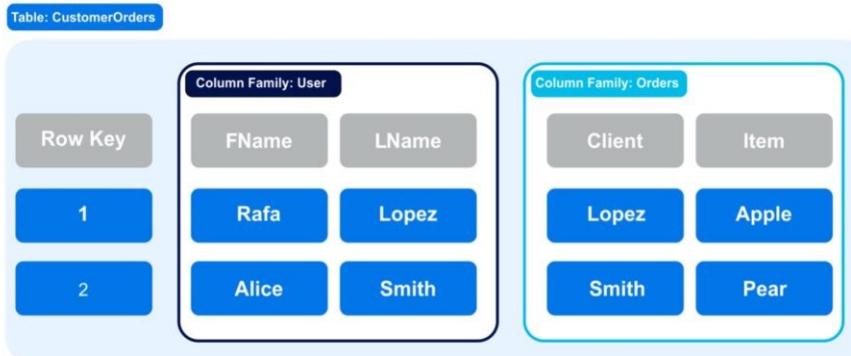
- Apache Cassandra, Apache HBase, Apache Accumulo, Hypertable, Google Bigtable

Column Databases – Approaches

❖ Apache HBase versus Apache Cassandra

– HBase *column*

- data model is the column-oriented table
- rows are divided into related columns of data called column families



```
> get 'CustomerOrders', '1'  
COLUMN          CELL  
User: FName timestamp = 1675123184293, value = Rafa  
User: LName timestamp = 1675123184293, value = Lopez  
Orders: Client timestamp = 1675123388213, value = Lopez  
Orders: Client timestamp = 1675123388214, value = Apple  
4 row(s) in 0.0260 seconds
```

– Cassandra: *row*

- data model is best described as a partitioned row store
- at top-level data model, keyspaces are column-families



```
cqlsh> select * from User;  
id | fname | lname  
---+---+---  
1 | Rafa | Lopez  
2 | Alice | Smith  
(2 rows)
```

Graph Databases – Basics

món: entidade

arestas: relação

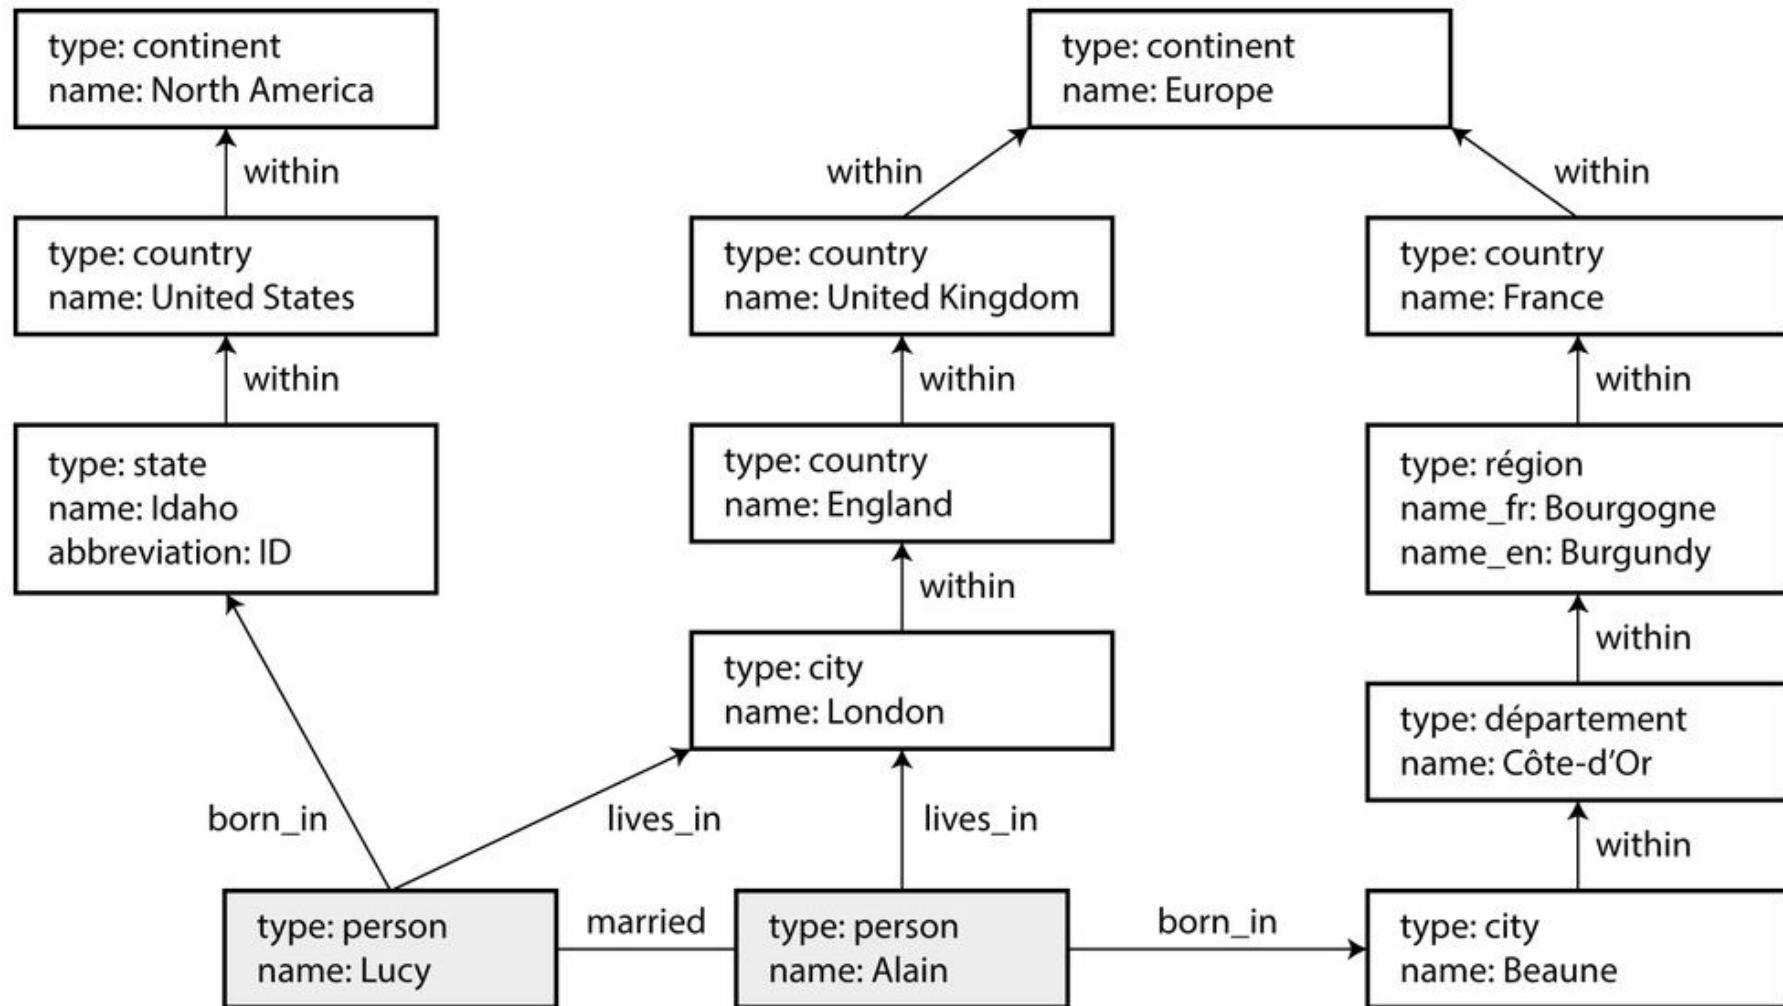
❖ Data Model

- Focus on modelling graphs' structure and properties
- Directed / undirected graphs, i.e. collections of ...
 - nodes (vertices) for real-world entities, and
 - relationships (edges) between these nodes
- Both the nodes and relationships can have properties

❖ Query patterns

- Create, update or remove a node / relationship in a graph
 - Graph algorithms (shortest paths, spanning trees, ...)
 - General graph traversals
 - Sub-graph queries or super-graph queries
 - Similarity based queries (approximate matching)

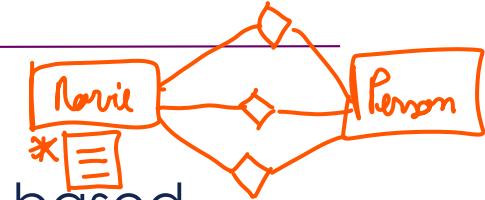
Graph Databases – Basics



Graph Databases – Basics

tom vértices e arestas
em que posso guardar nos nós
entendendo das mozes (os nó é uma
entrada*)!

permitem através de
algoritmos executar algoritmos
mais rápido e curto



❖ Suitable use cases

- Social networks, routing, dispatch, and location-based services, recommendation engines, chemical compounds, biological pathways, linguistic trees, ...

❖ When not to use

- Extensive batch operations are required
 - Multiple nodes / relationships are to be affected
- Too large graphs to be stored
 - Graph distribution is difficult or impossible at all

❖ Examples

- Neo4j, Titan, Apache Giraph, InfiniteGraph, FlockDB
- Multi-model: OrientDB, OpenLink Virtuoso, ArangoDB

Native XML Databases – Basics (JSON-like)

❖ Data model

- XML documents
- **Tree structure** with nested elements, attributes, and text values (beside other less important constructs)
- Documents are organized into collections

❖ Query languages

- **XPath**: XML Path Language (navigation)
- **XQuery**: XML Query Language (querying)
- **XSLT**: XSL Transformations (transformation)

❖ Examples

- Sedna, Tamin, BaseX, eXist-db
- Multi-model: MarkLogic, OpenLink Virtuoso

RDF Databases – Basics

(linguagem própria)

❖ Data model

- RDF **triples**
 - Components: **subject**, **predicate**, and **object**
 - Each triple represents a statement about a real-world entity
- **Triples can be viewed as graphs**
 - Vertices for subjects and objects
 - Edges directly correspond to individual statements

❖ Query language

- **SPARQL**: SPARQL Protocol and RDF Query Language

❖ Examples

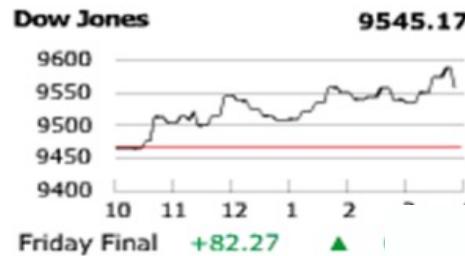
- Apache Jena, rdf4j (Sesame), Algebraix
- Multi-model: MarkLogic, OpenLink Virtuoso

Time Series Databases – basics

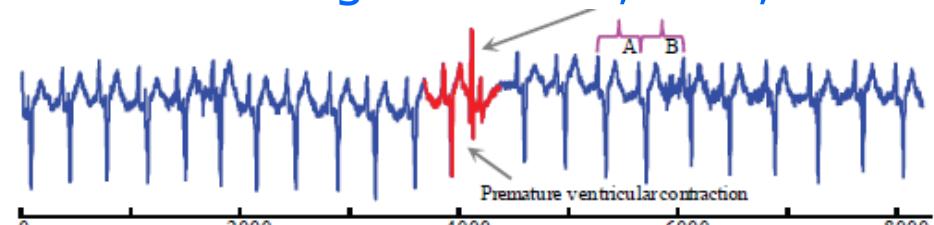
- ❖ Data model
 - Stores pairs “Time:Value”
- ❖ Query language
 - **Proprietary**: InfluxQL, ...
 - **SQL**: some multi-model engines
- ❖ Usage
 - store profiles, curves, traces or trends
 - fewer relationships between data entries
 - long sets of data
 - data patterns are “appreciated”
 - compression algorithms to manage the data efficiently
- ❖ Examples
 - InfluxDB, Prometheus, Graphite
 - Multi-model: Kdb, TimescaleDB

Time Series Examples

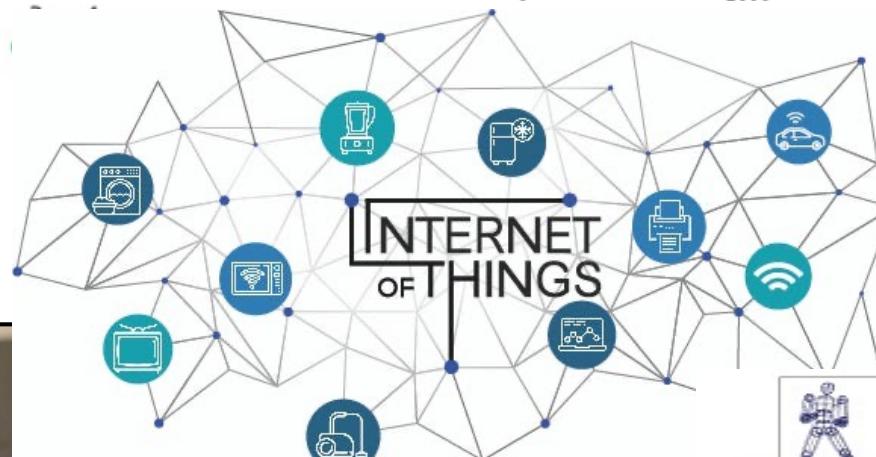
Stocks Data



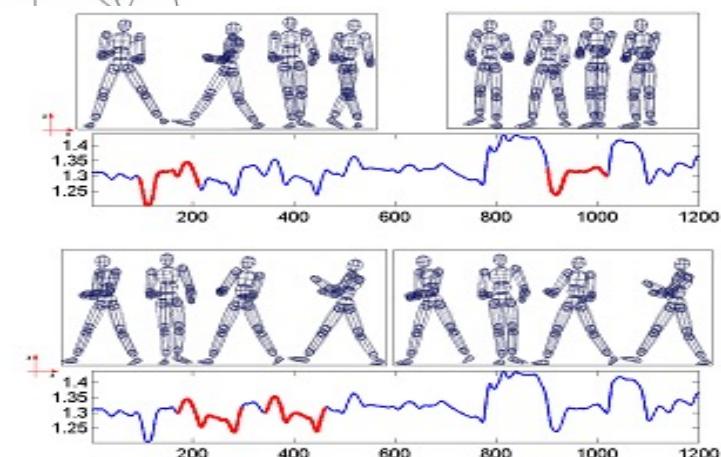
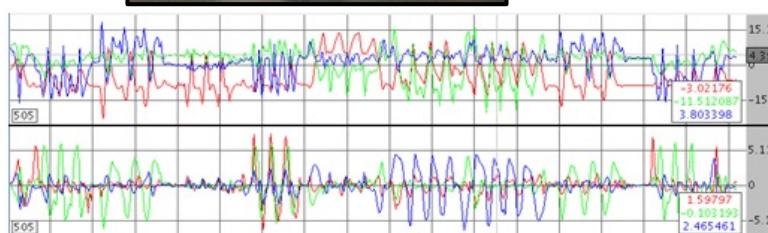
Vital Signals: ECG, EEG,...



Sensors



Motion Data



NoSQL Databases

- ❖ The end of relational databases?
- ❖ **Certainly no!**
 - They are still suitable for most projects (90%)
 - Familiarity, stability, feature set, available support, ...
- ❖ However, we should also consider different database models and systems
 - **Polyglot persistence** = usage of different data stores in different circumstances

Databases and data connectivity

❖ Relational model

❖ NoSQL models

- Key-value stores
- Document stores
- Column stores
- Graph databases

key-value model

column-family model

relational model

graph model

unrelated records

highly connected data

What next?

- ❖ Basic principles
 - Data formats: JSON, YAML, XML, RDF, ...
 - Distribution, scaling, sharding, replication, consistency
 - Parallelism, transactions, visualization, processing of graphs
- ❖ NoSQL technologies: principles, models, interfaces, languages, ...
 - Core databases: Redis, MongoDB, Cassandra, Neo4j
 - MapReduce: Apache Hadoop

Resources

- ❖ Martin Knopmann, **Designing Data-Intensive Applications**, O'Reilly Media, Inc., 2017.
- ❖ Pramod J Sadalage and Martin Fowler, **NoSQL Distilled** Addison-Wesley, 2016.
- ❖ Eric Redmond, Jim R. Wilson. **Seven databases in seven weeks**, Pragmatic Bookshelf, 2012.
- ❖ Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom, **Database systems: the complete book** (2nd Ed.), Pearson Education, 2009.